

Stroke (Brain Attack) Prediction



Context

According to the World Health Organization (WHO) stroke is the 2nd leading cause of death globally, responsible for approximately 11% of total deaths.

In recent years, strokes have emerged as a significant public health concern, impacting millions of individuals worldwide. A stroke, often referred to as a "brain attack," is a medical condition that occurs when the blood supply to the brain is interrupted or reduced, leading to severe health consequences and, in some cases, fatalities. Identifying individuals at high risk of stroke is crucial for early intervention and preventive measures.

This dataset is used to predict whether a patient is likely to get stroke based on the input parameters like gender, age, various diseases, and smoking status. Each row in the data provides relevant information about the patient.

Dataset Link: <https://www.kaggle.com/datasets/fedesoriano/stroke-prediction-dataset>
[\(https://www.kaggle.com/datasets/fedesoriano/stroke-prediction-dataset\)](https://www.kaggle.com/datasets/fedesoriano/stroke-prediction-dataset)

Attribute

1. id: unique identifier
2. gender: "Male", "Female" or "Other"
3. age: age of the patient
4. hypertension: 0 if the patient doesn't have hypertension, 1 if the patient has hypertension
5. heart_disease: 0 if the patient doesn't have any heart diseases, 1 if the patient has a heart disease
6. ever_married: "No" or "Yes"
7. work_type: "children", "Govt_jov", "Never_worked", "Private" or "Self-employed"
8. Residence_type: "Rural" or "Urban"
9. avg_glucose_level: average glucose level in blood
10. bmi: body mass index
11. smoking_status: "formerly smoked", "never smoked", "smokes" or "Unknown"*
12. stroke: 1 if the patient had a stroke or 0 if not

*Note: "Unknown" in smoking_status means that the information is unavailable for this patient

Some Collected Information/Facts From Sources

1. Stroke is the second most common cause of death in India. About 1,85,000 stroke cases are reported every year in India with nearly one stroke every 40 seconds. One stroke death every 4 min. (Source: AIIMS)
2. 5.2 millions strokes were witnessed in children, aged less than 20 years. More risk after 50's.
3. The chance of having a stroke about doubles every 10 years after age 55.
4. The cumulative incidence of stroke in India ranged from 105 to 152/100,000 persons per year.
5. Main Factors of having Stroke: Highblood Pressure, Diabetes, High Cholesterol and Obesity.
6. Cigarette smoking is a well established risk factor for all forms of stroke.
7. Glucose can result in inc. fatty deposits or clots in blood vessels.
8. Women generally live longer than men, more women have strokes over their lifetimes. Women also have unique risk factors for stroke, including: Having high blood pressure during pregnancy. Using certain types of birth control medicines, especially if they also smokes.

Aim

Our Stroke Prediction Project harnesses data science and machine learning algorithms to forecast stroke risk, allowing for early intervention and tailored care. Our goal is to improve patient outcomes and alleviate the strain on healthcare systems through precise predictive modeling.

Data Analysis

Import Libraries

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import warnings # Optional
warnings.filterwarnings('ignore')
```

```
In [4]: # create a dataframe from dataset
df = pd.read_csv('stroke data.csv')
```

```
In [5]: # check initial 5 records  
df.head()
```

```
Out[5]:
```

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status	stroke
0	9046	Male	67.0	0	1	Yes	Private	Urban	148.0	33.2	Never	0
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural	138.0	33.2	Ever	1
2	31112	Male	80.0	0	1	Yes	Private	Rural	148.0	33.2	Ever	1
3	60182	Female	49.0	0	0	Yes	Private	Urban	140.0	33.2	Never	0
4	1665	Female	79.0	1	0	Yes	Self-employed	Rural	140.0	33.2	Ever	1

```
In [6]: # information of dataframe  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 5110 entries, 0 to 5109  
Data columns (total 12 columns):  
 #   Column           Non-Null Count  Dtype     
---  --    
 0   id              5110 non-null    int64    
 1   gender          5110 non-null    object    
 2   age              5110 non-null    float64   
 3   hypertension     5110 non-null    int64    
 4   heart_disease   5110 non-null    int64    
 5   ever_married    5110 non-null    object    
 6   work_type        5110 non-null    object    
 7   Residence_type  5110 non-null    object    
 8   avg_glucose_level 5110 non-null    float64   
 9   bmi              4909 non-null    float64   
 10  smoking_status  5110 non-null    object    
 11  stroke           5110 non-null    int64    
dtypes: float64(3), int64(4), object(5)  
memory usage: 479.2+ KB
```

```
In [7]: # stats behind our dataframe  
df.describe()
```

Out[7]:

	id	age	hypertension	heart_disease	avg_glucose_level	bmi
count	5110.000000	5110.000000	5110.000000	5110.000000	5110.000000	4909.000000
mean	36517.829354	43.226614	0.097456	0.054012	106.147677	28.893237
std	21161.721625	22.612647	0.296607	0.226063	45.283560	7.854067
min	67.000000	0.080000	0.000000	0.000000	55.120000	10.300000
25%	17741.250000	25.000000	0.000000	0.000000	77.245000	23.500000
50%	36932.000000	45.000000	0.000000	0.000000	91.885000	28.100000
75%	54682.000000	61.000000	0.000000	0.000000	114.090000	33.100000
max	72940.000000	82.000000	1.000000	1.000000	271.740000	97.600000

```
In [8]: # correlation of dataframe between all columns (without Label encoding of column)  
df.corr()
```

Out[8]:

	id	age	hypertension	heart_disease	avg_glucose_level	bmi
id	1.000000	0.003538	0.003550	-0.001296	0.001092	0.003084
age	0.003538	1.000000	0.276398	0.263796	0.238171	0.333398
hypertension	0.003550	0.276398	1.000000	0.108306	0.174474	0.167811
heart_disease	-0.001296	0.263796	0.108306	1.000000	0.161857	0.041357
avg_glucose_level	0.001092	0.238171	0.174474	0.161857	1.000000	0.175502
bmi	0.003084	0.333398	0.167811	0.041357	0.175502	1.000000
stroke	0.006388	0.245257	0.127904	0.134914	0.131945	0.042374

```
In [9]: # stats behind our dataframe including object datatype  
df.describe(include='object')
```

Out[9]:

	gender	ever_married	work_type	Residence_type	smoking_status
count	5110	5110	5110	5110	5110
unique	3	2	5	2	4
top	Female	Yes	Private	Urban	never smoked
freq	2994	3353	2925	2596	1892

```
In [10]: # check the unique value in object datatype column  
df['gender'].value_counts()
```

```
Out[10]: Female    2994  
Male      2115  
Other       1  
Name: gender, dtype: int64
```

```
In [11]: df['work_type'].value_counts()
```

```
Out[11]: Private     2925  
Self-employed    819  
children        687  
Govt_job         657  
Never_worked      22  
Name: work_type, dtype: int64
```

```
In [12]: df['smoking_status'].value_counts()
```

```
Out[12]: never smoked    1892  
Unknown          1544  
formerly smoked   885  
smokes           789  
Name: smoking_status, dtype: int64
```

```
In [13]: # check the null values in dataframe  
df.isna().sum()
```

```
Out[13]: id            0  
gender          0  
age             0  
hypertension    0  
heart_disease   0  
ever_married    0  
work_type        0  
Residence_type   0  
avg_glucose_level 0  
bmi             201  
smoking_status   0  
stroke           0  
dtype: int64
```

Data Preprocessing 1

```
In [14]: '''Remove 'id' columns bcz this is not really useful for our model, and if we  
don't remove it our model consider it as real value and make wrong predictions  
  
df.drop('id',axis=1,inplace=True)
```

```
In [15]: df.head()
```

```
Out[15]:
```

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_gluc
0	Male	67.0	0	1	Yes	Private	Urban	30.0
1	Female	61.0	0	0	Yes	Self-employed	Rural	24.0
2	Male	80.0	0	1	Yes	Private	Rural	33.6
3	Female	49.0	0	0	Yes	Private	Urban	23.3
4	Female	79.0	1	0	Yes	Self-employed	Rural	34.6



```
In [16]: # Remove null values in 'bmi' column  
df.dropna(axis=0,inplace=True)
```

```
In [17]: # identifying index of 'Other' value in 'Gender' column  
df[df['gender']=='Other']
```

```
Out[17]:
```

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_gluc
3116	Other	26.0	0	0	No	Private	Rural	30.0



```
In [18]: # Remove 'Other' value from gender column bcz their is only value in that column  
df = df.drop([3116],axis=0)
```

```
In [19]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 4908 entries, 0 to 5109  
Data columns (total 11 columns):  
 #   Column           Non-Null Count  Dtype     
---  --  
 0   gender          4908 non-null    object    
 1   age              4908 non-null    float64  
 2   hypertension     4908 non-null    int64     
 3   heart_disease   4908 non-null    int64     
 4   ever_married    4908 non-null    object    
 5   work_type        4908 non-null    object    
 6   Residence_type  4908 non-null    object    
 7   avg_glucose_level 4908 non-null    float64  
 8   bmi              4908 non-null    float64  
 9   smoking_status   4908 non-null    object    
 10  stroke           4908 non-null    int64     
dtypes: float64(3), int64(3), object(5)  
memory usage: 460.1+ KB
```

```
In [20]: # exploring dataframe using group by of 'work type' column
work_type_gb = df.groupby('work_type')
work_type_gb.mean()
```

Out[20]:

	age	hypertension	heart_disease	avg_glucose_level	bmi	stroke
work_type						
Govt_job	50.717460	0.106349	0.052381	107.356825	30.522063	0.044444
Never_worked	16.181818	0.000000	0.000000	96.042727	25.545455	0.000000
Private	45.210676	0.090747	0.049466	105.647591	30.307438	0.045196
Self-employed	59.916129	0.166452	0.090323	112.389161	30.211871	0.068387
children	6.877973	0.000000	0.001490	94.009806	20.038003	0.001490

```
In [21]: # exploring dataframe using group by of 'Residence type' column
residence_type_gb = df.groupby('Residence_type')
residence_type_gb.mean()
```

Out[21]:

	age	hypertension	heart_disease	avg_glucose_level	bmi	stroke
Residence_type						
Rural	42.621737	0.092225	0.050041	105.632821	28.896898	0.041356
Urban	43.108739	0.091566	0.048996	104.971683	28.892289	0.043775

```
In [22]: # exploring dataframe using group by of 'smoking_status' column
smoking_status_gb = df.groupby('smoking_status')
smoking_status_gb.mean()
```

Out[22]:

	age	hypertension	heart_disease	avg_glucose_level	bmi	stroke
smoking_status						
Unknown	29.511207	0.028995	0.024949	98.335927	25.666352	0.019555
formerly smoked	54.958134	0.131579	0.083732	112.543038	30.757177	0.068182
never smoked	46.469222	0.116631	0.043737	107.136215	29.982559	0.045356
smokes	46.986431	0.111262	0.074627	106.465699	30.543555	0.052917

```
In [23]: # exploring dataframe using group by of 'gender' column
gender_gb = df.groupby('gender')
gender_gb.mean()
```

Out[23]:

	age	hypertension	heart_disease	avg_glucose_level	bmi	stroke
gender						
Female	43.437805	0.086641	0.034518	103.329914	29.065758	0.041422
Male	42.049130	0.099453	0.071109	108.131721	28.647936	0.044257

```
In [24]: # exploring dataframe using group by of 'ever_married' column  
ever_married_gb = df.groupby('ever_married')  
ever_married_gb.mean()
```


```
Out[24]:
```

ever_married	age	hypertension	heart_disease	avg_glucose_level	bmi	stroke
No	21.815798	0.027582	0.016432	96.059742	25.216373	0.013498
Yes	54.065543	0.126092	0.067104	110.210315	30.850749	0.058052

Exploratory Data Analysis (EDA)

```
In [25]: # let's create a lazy plot called pairplot for normal analysis of complete data  
sns.pairplot(df,hue='stroke')
```


```
Out[25]: <seaborn.axisgrid.PairGrid at 0x7cf82963550>
```



Univariate Analysis


```
In [26]: # check the distribution of dataset features  
df.hist(bins=60,figsize=(15,10))  
plt.suptitle('Feature Distribution',x=0.5,y=1.02,ha='center',fontsize='large')
```

Out[26]: Text(0.5, 1.02, 'Feature Distribution')



```
In [27]: # check the distribution of dataset features using barplot
df.boxplot(figsize=(12,8))
plt.suptitle('Feature Distribution',x=0.5,y=1.02,ha='center',fontsize='large')
```

Out[27]: Text(0.5, 1.02, 'Feature Distribution')



As we are able to see that column 'avg glucose level' and 'bmi' has more outliers as compare to others.


```
In [28]: # List for the gender column
gender = ['Female', 'Male']

# make an object which store value count of the elements in gender column
gender_data = pd.Series(df['gender']).value_counts()

# adjust size of the pie chart
plt.figure(figsize=(4,4))

# plotting pie chart
plt.pie(gender_data, labels=gender, autopct='%.0f%%')
```

```
Out[28]: ([<matplotlib.patches.Wedge at 0x7cf7ae8fd30>,
<matplotlib.patches.Wedge at 0x7cf7ae8fc40>],
[Text(-0.307755600196639, 1.0560712525902818, 'Female'),
Text(0.30775560019663867, -1.0560712525902818, 'Male')],
[Text(-0.16786669101634855, 0.5760388650492445, '59%'),
Text(0.16786669101634835, -0.5760388650492446, '41%')])
```




```
In [29]: # List for the married column
married = ['Married','Not Married']

# make an object which store value count of the elements in married column
married_data = pd.Series(df['ever_married']).value_counts()

# adjust size of the pie chart
plt.figure(figsize=(4,4))

# ploting pie chart
plt.pie(married_data,labels=married,autopct='%.0f%%')
```

```
Out[29]: ([<matplotlib.patches.Wedge at 0x7cfa820d8eb0>,
<matplotlib.patches.Wedge at 0x7cfa820d8820>],
[Text(-0.5080276543906456, 0.9756576768387256, 'Married'),
Text(0.5080277457383074, -0.9756576292737397, 'Not Married')],
[Text(-0.27710599330398844, 0.5321769146393048, '65%'),
Text(0.2771060431299858, -0.532176888694767, '35%')])
```




```
In [30]: # List for the work_type column
work_type = ['Private', 'Self-Employed', 'Children', 'Govt. Job', 'Never Worked']

# make an object which store value count of the elements in work_type column
work_type_data = pd.Series(df['work_type']).value_counts()

# adjust size of the pie chart
plt.figure(figsize=(6,4))

# plotting pie chart
plt.pie(work_type_data, labels=work_type, autopct='%.0f%%')
```

```
Out[30]: ([<matplotlib.patches.Wedge at 0x7cfa7cfcb1f0>,
<matplotlib.patches.Wedge at 0x7cfa7cfcb310>,
<matplotlib.patches.Wedge at 0x7cfa7d022470>,
<matplotlib.patches.Wedge at 0x7cfa7d023f10>,
<matplotlib.patches.Wedge at 0x7cfa7cfe5690>],
[Text(-0.24849794293519883, 1.0715637042924582, 'Private'),
Text(-0.6382191091776186, -0.8959220773485421, 'Self-Employed'),
Text(0.33200713375555313, -1.0486997964791556, 'Children'),
Text(0.9992078303999088, -0.4599822949500419, 'Govt. Job'),
Text(1.0998909305275677, -0.015490027217579071, 'Never Worked'),
[Text(-0.13554433251010845, 0.5844892932504316, '57%'),
Text(-0.3481195140968828, -0.4886847694628411, '16%'),
Text(0.1810948002303017, -0.5720180708068121, '14%'),
Text(0.5450224529454047, -0.25089943360911376, '13%'),
Text(0.5999405075604914, -0.008449105755043129, '0%')])
```




```
In [31]: # List for the Residence_type column
Residence_type = ['Urban', 'Rural']

# make an object which store value count of the elements in Residence_type col
residence_type_data = pd.Series(df['Residence_type']).value_counts()

# adjust size of the pie chart
plt.figure(figsize=(6,4))

# plotting pie chart
plt.pie(residence_type_data, labels=Residence_type, autopct='%.0f%%')
```

```
Out[31]: ([<matplotlib.patches.Wedge at 0x7cf7aee3400>,
<matplotlib.patches.Wedge at 0x7cf7aee3ca0>],
[Text(-0.025345515982636815, 1.099707963424642, 'Urban'),
Text(0.025345413020546048, -1.0997079657976558, 'Rural')],
[Text(-0.01382482689962008, 0.599840707322532, '51%'),
Text(0.013824770738479661, -0.5998407086169031, '49%')])
```




```
In [32]: # List for the smoking_status column
smoking_status = ['Never Smoked', 'Unknown', 'Formerly Smoked', 'Smokes']

# make an object which store value count of the elements in smoking_status column
smoking_status_data = pd.Series(df['smoking_status']).value_counts()

# adjust size of the pie chart
plt.figure(figsize=(6,4))

# plotting pie chart
plt.pie(smoking_status_data, labels=smoking_status, autopct='%.0f%%')
```

```
Out[32]: ([<matplotlib.patches.Wedge at 0x7cf86a620>,
<matplotlib.patches.Wedge at 0x7cf86bf10>,
<matplotlib.patches.Wedge at 0x7cf82107940>,
<matplotlib.patches.Wedge at 0x7cf87ffd3790>],
[Text(0.41345956848245324, 1.019338601854312, 'Never Smoked'),
Text(-1.0825052746740849, -0.1954029946105854, 'Unknown'),
Text(0.10124768212643617, -1.0953305012022738, 'Formerly Smoked'),
Text(0.97985128217811, -0.49989145303146976, 'Smokes')],
[Text(0.225523400990429, 0.5560028737387156, '38%'),
Text(-0.5904574225495007, -0.10658345160577384, '30%'),
Text(0.055226008432601545, -0.5974530006557857, '17%'),
Text(0.5344643357335145, -0.27266806528989257, '15%')])
```




```
In [33]: # Let's see the distribution of real numbers columns with the help of kde plot
plt.figure(figsize=(15,15))

plt.subplot(3,1,1)
sns.distplot(x=df['age'],kde=True)
plt.title('Age')

plt.subplot(3,1,2)
sns.distplot(x=df['avg_glucose_level'],kde=True)
plt.title('Glucose Level')


plt.subplot(3,1,3)
sns.distplot(x=df['bmi'],kde=True)
plt.title('Body Mass Index')

plt.tight_layout()
```




```
In [34]: # create a countplot for Gender column  
plt.figure(figsize=(10,6))  
sns.countplot(x='gender',data=df)
```

Out[34]: <Axes: xlabel='gender', ylabel='count'>




```
In [35]: # create a countplot for Work Type column  
plt.figure(figsize=(10,6))  
sns.countplot(x='work_type',data=df)
```

Out[35]: <Axes: xlabel='work_type', ylabel='count'>




```
In [36]: # create a countplot for Smoking Status column  
plt.figure(figsize=(10,6))  
sns.countplot(x='smoking_status',data=df)
```

```
Out[36]: <Axes: xlabel='smoking_status', ylabel='count'>
```



```
In [37]: # create a countplot for target variable 'stroke' column  
sns.countplot(x='stroke',data=df)
```

```
Out[37]: <Axes: xlabel='stroke', ylabel='count'>
```



Note: Here we are able to see that provided dataset is imbalance means in our dataset most number of patients are not suffering from stroke (4500+) and minimum number of patients suffer from stroke around (200 - 300). So, this will effect on our machine learning model future. Hence, we will just continue this project for learning purpose, although our model perform worse for real time data and it will perform wrong predictions.

0 indicates patient not suffer from stroke and 1 indicates patient suffer from stroke.

Bivariate Analysis


```
In [38]: # Let's plot the boxplot for all numeric continuous columns w.r.t target feature
plt.figure(figsize=(15,15))

plt.subplot(3,1,1)
sns.boxplot(x=df['stroke'],y=df['age'],data=df)
plt.title('Age')

plt.subplot(3,1,2)
sns.boxplot(x=df['stroke'],y=df['avg_glucose_level'],data=df)
plt.title('Glucose Level')

plt.subplot(3,1,3)
sns.boxplot(x=df['stroke'],y=df['bmi'],data=df)
plt.title('Body Mass Index')

plt.tight_layout()
```



```
In [39]: # checking all the features correlation w.r.t target variable

# adjust the size of the plot
plt.figure(figsize=(12,5))


# plot correlation
df.corr()['stroke'][:-1].plot(kind='bar')

# create a 'x Label'
plt.xlabel('Features')

# create a 'y Label'
plt.ylabel('Correlation Value')

# create a title for the plot
plt.title('Correlation Plot of All Features')
```

Out[39]: Text(0.5, 1.0, 'Correlation Plot of All Features')




```
In [40]: # Let's plot the violinplot for all numeric continuous columns w.r.t target fea
plt.figure(figsize=(15,15))

plt.subplot(3,1,1)
sns.violinplot(x=df['stroke'],y=df['age'],data=df)
plt.title('Age')

plt.subplot(3,1,2)
sns.violinplot(x=df['stroke'],y=df['avg_glucose_level'],data=df)
plt.title('Glucose Level')

plt.subplot(3,1,3)
sns.violinplot(x=df['stroke'],y=df['bmi'],data=df)
plt.title('Body Mass Index')

plt.tight_layout()
```




```
In [41]: # check the relation between 'Age' vs 'Glucose Level' and 'Age' vs 'BMI' with
plt.figure(figsize=(15,15))

plt.subplot(3,1,1)
sns.scatterplot(x=df['age'],y=df['avg_glucose_level'],hue='stroke',data=df)
plt.title('Age vs Glucose Level')


plt.subplot(3,1,2)
sns.scatterplot(x=df['age'],y=df['bmi'],hue='stroke',data=df)
plt.title('Age vs BMI')

plt.tight_layout()
```



```
In [42]: # check the relation between 'BMI' vs 'Glucose Level' with scatterplot  
plt.figure(figsize=(15,10))  
  
sns.scatterplot(x=df['bmi'],y=df['avg_glucose_level'],hue='stroke',data=df)  
  
plt.title('BMI vs Glucose Level')
```

Out[42]: Text(0.5, 1.0, 'BMI vs Glucose Level')



Multi Variate Analysis


```
In [43]: # create a heatmap to visualize the correlation between all features

# create a variable for correlation
df_corr = df.corr()

# plot heatmap
sns.heatmap(df_corr, annot=True, linecolor='white', linewidths=0.2)

# set a title
plt.title('Heatmap of Dataframe (Without Label Encoding)')
```

Out[43]: Text(0.5, 1.0, 'Heatmap of Dataframe (Without Label Encoding)')



Data Preprocessing 2 (Label Encoding)

```
In [44]: # import sklearn function 'LabelEncoder' for converting categorical values into numbers
from sklearn.preprocessing import LabelEncoder

# make a variable for LabelEncoder function
label_encoder = LabelEncoder()
```

'Gender' column Label Encoding

```
In [45]: # Label encoding of 'Gender' column in dataframe
df['gender'] = label_encoder.fit_transform(df['gender'])
```

'Ever Married' column Label Encoding

```
In [46]: # Label encoding of 'Ever Married' column in dataframe
df['ever_married'] = label_encoder.fit_transform(df['ever_married'])
```

'Work Type' column Label Encoding

```
In [47]: # Label encoding of 'Work Type' column in dataframe
df['work_type'] = label_encoder.fit_transform(df['work_type'])
```

'Residence Type' column Label Encoding

```
In [48]: # Label encoding of 'Residence Type' column in dataframe
df['Residence_type'] = label_encoder.fit_transform(df['Residence_type'])
```

'Smoking Status' column Label Encoding

```
In [49]: # Label encoding of 'Smoking Status' column in dataframe
df['smoking_status'] = label_encoder.fit_transform(df['smoking_status'])
```

```
In [50]: # Let's us take a look that our features encoded or not  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 4908 entries, 0 to 5109  
Data columns (total 11 columns):  
 #   Column           Non-Null Count  Dtype     
---  --    
 0   gender          4908 non-null    int64    
 1   age              4908 non-null    float64    
 2   hypertension     4908 non-null    int64    
 3   heart_disease    4908 non-null    int64    
 4   ever_married     4908 non-null    int64    
 5   work_type        4908 non-null    int64    
 6   Residence_type   4908 non-null    int64    
 7   avg_glucose_level 4908 non-null    float64    
 8   bmi              4908 non-null    float64    
 9   smoking_status   4908 non-null    int64    
 10  stroke           4908 non-null    int64    
dtypes: float64(3), int64(8)  
memory usage: 460.1 KB
```

Boom! As we are able to see that our all features now in numeric values without increasing the number of columns in dataframe, we can use get_dummies function of pandas library but that increase the number of features in dataframe that will effect on our machine learning model.

Heatmap

```
In [51]: # create a heatmap to visualize the correlation between all features


# adjust the size of heatmap
plt.figure(figsize=(15,10))

# create a variable for correlation
df_corr = df.corr()

# plot heatmap
sns.heatmap(df_corr, annot=True, linecolor='white', linewidths=0.2)

# set a title
plt.title('Heatmap of Dataframe (After Label Encoding)')
```

Out[51]: Text(0.5, 1.0, 'Heatmap of Dataframe (After Label Encoding)')



```
In [52]: # checking all the features correlation w.r.t target variable

# adjust the size of the plot
plt.figure(figsize=(12,5))


# plot correlation
df.corr()['stroke'][:-1].sort_values().plot(kind='bar')

# create a 'x Label'
plt.xlabel('Features')

# create a 'y Label'
plt.ylabel('Correlation Value')

# create a title for the plot
plt.title('Correlation Plot of All Features')
```

Out[52]: Text(0.5, 1.0, 'Correlation Plot of All Features')



Machine Learning Model

Import All Libaries Used To Build Model

```
In [54]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import classification_report,confusion_matrix
from sklearn.metrics import accuracy_score,log_loss
from sklearn.metrics import roc_auc_score,roc_curve
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.naive_bayes import MultinomialNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import GridSearchCV
import shap
```

Splitting of Data

```
In [55]: # splitting data into dependent and independent variables


# independent variable
X = df.drop('stroke',axis=1)

# dependent variable
y = df['stroke']
```

Sampling Data (Handling Imbalanced Dataset)

```
In [56]: # check the original dataset classes  
sns.countplot(x=y,data=df)
```

```
Out[56]: <Axes: xlabel='stroke', ylabel='count'>
```




```
In [57]: # using sampling technique to balance the classes
```

```
# import library  
from imblearn.over_sampling import RandomOverSampler  
from collections import Counter  
ros = RandomOverSampler(random_state=42)  
  
X_res, y_res = ros.fit_resample(X, y)  
  
print('Original dataset shape', Counter(y))  
print('Resample dataset shape', Counter(y_res))
```

```
Original dataset shape Counter({0: 4699, 1: 209})  
Resample dataset shape Counter({1: 4699, 0: 4699})
```

```
In [58]: # check the dataset classes after resampling  
sns.countplot(x=y_res,data=df)
```

```
Out[58]: <Axes: xlabel='stroke', ylabel='count'>
```



```
In [59]: # splitting data into training and test data  
X_train, X_test, y_train, y_test = train_test_split(X_res, y_res, test_size=0.2)
```

```
In [60]: # shape of all training and testing data variables  
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
Out[60]: ((6578, 10), (2820, 10), (6578,), (2820,))
```

Scaling Train and Test Data

```
In [61]: # make a variable for 'MinMaxScale' function  
scaler = MinMaxScaler()  
  
# scaling training data variable  
X_train = scaler.fit_transform(X_train)  
  
# scaling testing data variable  
X_test = scaler.transform(X_test)
```

Making an Object

```
In [62]: # making an object for all machine-learning algorithms
lr = LogisticRegression()
dtree = DecisionTreeClassifier()
rfc = RandomForestClassifier(n_estimators=200)
svm = SVC()
knn = KNeighborsClassifier(n_neighbors=5)
nb = MultinomialNB()
```

Accuracy Check For All Algorithms

```
In [63]: # making a List of all used ml algorithms
ml_names = ['Logistic Reg.', 'Decision Tree', 'Random Forest', 'SVM', 'K Nearest N
             'Naive Bayes']

# making a List of all used ml algorithms object
ml_object = [lr, dtree, rfc, svm, knn, nb]

# using for Loop to check the accuracy for all algorithms
print('Accuracy of all Algorithms\n')
# using both lists in our loop
for i, j in zip(ml_names, ml_object):

    # training our all algorithms
    j.fit(X_train, y_train)

    # make predictions for all algorithms with the help of test data
    pred = j.predict(X_test)

    # printing accuracy for all algorithms
    print(f'{i} : {accuracy_score(y_test, pred)*100:.2f}\n')
```

Accuracy of all Algorithms

Logistic Reg. : 78.05%

Decision Tree : 97.59%

Random Forest : 99.40%

SVM : 81.60%

K Nearest Neighbor : 93.44%

Naive Bayes : 66.49%

Log Loss Check for All Algorithms

```
In [64]: # making a list of all used ml algorithms
ml_names = ['Logistic Reg.', 'Decision Tree', 'Random Forest', 'SVM', 'K Nearest N
             'Naive Bayes']

# making a list of all used ml algorithms object
ml_object = [lr,dtree,rfc,svm,knn,nb]

# using for loop to check the Log Loss for all algorithms
print('Log Loss of all Algorithms\n')
# using both lists in our loop
for i,j in zip(ml_names,ml_object):

    # training our all algorithms
    j.fit(X_train,y_train)

    # make predictions for all algorithms with the help of test data
    pred = j.predict(X_test)

    # printing log loss for all algorithms
    print(f'{i} : {log_loss(y_test,pred):.2f}\n')
```

Log Loss of all Algorithms

Logistic Reg. : 7.91

Decision Tree : 0.92

Random Forest : 0.19

SVM : 6.63

K Nearest Neighbor : 2.36

Naive Bayes : 12.08

As we see all the algorithms works well and all gives 90%+ accuracy for our testing data. But, we need only one algorithms to make prediction. So, we are going to work on top 3 algorithms which work amazing above on the basis of Accuracy score (need max. score) and Log Loss score (need min. score).

Algorithms	Accurcay Score	Log Loss Score
1. Random Forest	99.40%	0.19
2. Decision Tree	97.59%	0.92

Evaluation of Selected Algorithms

```
In [65]: # making a List of selected ml algorithms
ml_names = ['Random Forest','Decision Tree','K Nearest Neighor']

# making a List of selected ml algorithms object
ml_object = [rfc,dtree,knn]

# using for Loop to check the confusion matric and classification report for s
# using both lists in our loop
for i,j in zip(ml_names,ml_object):
    print(f'{i} Evaluation:\n')
    # training our all algorithms
    j.fit(X_train,y_train)

    # make predictions for all algorithms with the help of test data
    pred = j.predict(X_test)

    # printing classification report for all algorithms
    print(f'{i} Classification Report:')
    print(f'{classification_report(y_test,pred)}\n')

    # printing confusion matrix for all algorithms
    print(f'{i} Confusion Matrix:')
    print(f'{confusion_matrix(y_test,pred)}\n')
    print('='*70)
```

Random Forest Evaluation:

Random Forest Classification Report:

	precision	recall	f1-score	support
0	1.00	0.99	0.99	1407
1	0.99	1.00	1.00	1413
accuracy			1.00	2820
macro avg	1.00	1.00	1.00	2820
weighted avg	1.00	1.00	1.00	2820

Random Forest Confusion Matrix:

```
[[1393  14]
 [  0 1413]]
```

=====

Decision Tree Evaluation:

Decision Tree Classification Report:

	precision	recall	f1-score	support
0	1.00	0.95	0.98	1407
1	0.95	1.00	0.98	1413
accuracy			0.98	2820
macro avg	0.98	0.98	0.98	2820
weighted avg	0.98	0.98	0.98	2820

Decision Tree Confusion Matrix:

```
[[1339  68]
 [  0 1413]]
```

=====

K Nearest Neighbor Evaluation:

K Nearest Neighbor Classification Report:

	precision	recall	f1-score	support
0	1.00	0.87	0.93	1407
1	0.88	1.00	0.94	1413
accuracy			0.93	2820
macro avg	0.94	0.93	0.93	2820
weighted avg	0.94	0.93	0.93	2820

K Nearest Neighbor Confusion Matrix:

```
[[1222 185]
 [  0 1413]]
```

=====

Random Forest

```
In [66]: # Performing Hyperparameter tuning for more accurate performance

# Define a dictionary of hyperparameters and their possible values
param_grid = {
    'n_estimators': [100, 200, 300],                      # Number of trees in the forest
    'max_depth': [10, 20, 30],                            # Maximum depth of individual t
    'min_samples_split': [2, 5, 10],                      # Minimum samples required to sp
    'min_samples_leaf': [1, 2, 4]                         # Minimum samples required to be
}

# Create a GridSearchCV object
grid_search = GridSearchCV(estimator=rfc, param_grid=param_grid, cv=5, scoring

# Fit the grid search to your training data
grid_search.fit(X_train, y_train)

# Get the best hyperparameters
best_params = grid_search.best_params_
print("Best Hyperparameters:", best_params)

# Get the best model
best_rfc_model = grid_search.best_estimator_

# Make predictions with the best model
rfc_pred = best_rfc_model.predict(X_test)
```


```
Best Hyperparameters: {'max_depth': 30, 'min_samples_leaf': 1, 'min_samples_s
plit': 2, 'n_estimators': 200}
```

```
In [67]: # making confusion matrix
cm = confusion_matrix(y_test, rfc_pred)


# adjust the size of the confusion matrix
plt.figure(figsize=(6, 6))
sns.set(font_scale=1.2)

# create a heatmap which show the confusion matrix of the random forest model
sns.heatmap(cm, annot=True, fmt='d', cmap='coolwarm', square=True, xticklabels=
plt.xlabel('Predicted Values') # x label of the confusion matrix
plt.ylabel('True Values') # y label of the confusion matrix
plt.title('Confusion Matrix for Random Forest') # title of the confusion matrix
```


Out[67]: Text(0.5, 1.0, 'Confusion Matrix for Random Forest')



```
In [68]: # shap plot for the random forest model
explainer = shap.TreeExplainer(rfc)
shap_values = explainer.shap_values(X_test)
shap.summary_plot(shap_values,X_test,feature_names=X.columns)
```



```
In [69]: explainer = shap.TreeExplainer(rfc)
shap_values = explainer.shap_values(X_test)
shap.summary_plot(shap_values[1],X_test,feature_names=X.columns)
```




```
In [70]: # ROC and AUC score for the random forest model
rfc_pred_prob = rfc.predict_proba(X_test)[:, :, 1]

rfc_actual_predict = pd.concat([pd.DataFrame(np.array(y_test), columns=['y actual'],
                                             index=y_test.index),
                                 pd.DataFrame(rfc_pred_prob, columns=['y pred probability'],
                                             index=y_test.index)], axis=1)

fpr, tpr, tr = roc_curve(rfc_actual_predict['y actual'], rfc_actual_predict['y pred probability'])
auc = roc_auc_score(rfc_actual_predict['y actual'], rfc_actual_predict['y pred probability'])

plt.plot(fpr, tpr, label='AUC=%f' % auc)
plt.plot(fpr, fpr, linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Random Forest')
plt.legend()
```

Out[70]: <matplotlib.legend.Legend at 0x7cf9c5865ea0>



Decision Tree

```
In [71]: # Define a dictionary of hyperparameters and their possible values
param_grid = {
    'criterion': ['gini', 'entropy'], # Splitting criterion
    'max_depth': [10, 20, 30],       # Maximum depth of the tree
    'min_samples_split': [2, 5, 10], # Minimum samples required to split an internal node
    'min_samples_leaf': [1, 2, 4]   # Minimum samples required to be at a leaf node
}

# Create a GridSearchCV object
grid_search = GridSearchCV(estimator=dtree, param_grid=param_grid, cv=5, scoring='accuracy')

# Fit the grid search to your training data
grid_search.fit(X_train, y_train)

# Get the best hyperparameters
best_params = grid_search.best_params_
print("Best Hyperparameters:", best_params)

# Get the best model
best_dtree_model = grid_search.best_estimator_

# Make predictions with the best model
dtree_pred = best_dtree_model.predict(X_test)
```

Best Hyperparameters: {'criterion': 'entropy', 'max_depth': 30, 'min_samples_leaf': 1, 'min_samples_split': 5}


```
In [72]: # making confusion matrix
cm = confusion_matrix(y_test,dtree_pred)

# adjust the size of the confusion matrix
plt.figure(figsize=(6, 6))
sns.set(font_scale=1.2)


# create a heatmap which show the confusion matrix of the decision tree model
sns.heatmap(cm, annot=True, fmt='d', cmap='coolwarm', square=True,xticklabels=
            yticklabels=['0', '1'])

plt.xlabel('Predicted Values') # x Label of the confusion matrix
plt.ylabel('True Values') # y Label of the confusion matrix
plt.title('Confusion Matrix for Decision Tree') # title of the confusion matrix
```


Out[72]: Text(0.5, 1.0, 'Confusion Matrix for Decision Tree')



```
In [73]: # shap plot for the decision tree model
explainer = shap.TreeExplainer(dtree)
shap_values = explainer.shap_values(X_test)
shap.summary_plot(shap_values,X_test,feature_names=X.columns)
```



```
In [74]: explainer = shap.TreeExplainer(dtree)
shap_values = explainer.shap_values(X_test)
shap.summary_plot(shap_values[1],X_test,feature_names=X.columns)
```




```
In [75]: # ROC and AUC score for the decision tree model
dtree_pred_prob = dtree.predict_proba(X_test)[:, :, 1]

dtree_actual_predict = pd.concat([pd.DataFrame(np.array(y_test), columns=['y actual']),
                                   pd.DataFrame(dtree_pred_prob, columns=['y pred probability'])],
                                 axis=1)
dtree_actual_predict.index = y_test.index

fpr, tpr, tr = roc_curve(dtree_actual_predict['y actual'], dtree_actual_predict['y pred probability'])
auc = roc_auc_score(dtree_actual_predict['y actual'], dtree_actual_predict['y pred probability'])

plt.plot(fpr, tpr, label='AUC=%0.4f' % auc)
plt.plot(fpr, fpr, linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Decision Tree')
plt.legend()
```

```
Out[75]: <matplotlib.legend.Legend at 0x7cf9c47332e0>
```



Saving Decision Tree Model

```
In [76]: import pickle
```

```
In [88]: # use dump() function to save the model with pickle  
save_dtrees_model = pickle.dumps(best_dtrees_model)  
  
# use Load() function to loading the saved model  
dtrees_from_pickle = pickle.loads(save_rfc_model)  
  
# after loading, use test data to predict the outcome  
dtrees_from_pickle.predict(X_test)
```

```
Out[88]: array([0, 0, 0, ..., 0, 1, 1])
```

```
In [89]: df.head()
```

```
Out[89]:   gender  age  hypertension  heart_disease  ever_married  work_type  Residence_type  avg_gluc  
0         1  67.0           0            1             1          2              1  
1         1  80.0           0            1             1          2              0  
2         0  49.0           0            0             1          2              1  
3         0  79.0           1            0             1          3              0  
4         1  81.0           0            0             1          2              1
```



```
In [90]: rfc_from_pickle.predict([[1,67.0,0,1,1,2,1,228.69,36.6,1]])
```

```
Out[90]: array([1])
```

Conclusions/Report

Data Analysis:

1. Working People's like: Govt. Job, Private and Self-Employed are has high 'Glucose Level' and 'BMI'.
2. Paitents who smoke formerly has high 'Glucose Level' and 'BMI'.
3. People who married (mean age 54 years), they suffer from high hypertension, heart disease, glucose level and BMI, that why they suffer from stroke.
4. Most of Stroke cases found in Females as compare to Males.
5. Most of the paitents Never Smoke but, still suffer from Stroke.
6. Mean BMI is around 29.5 approx.
7. **Note:** We have most number paitents data who not suffer from Stroke. This will impact on ML model. (Imbalance Dataset)
8. Feature 'Age' and 'Ever Married' are more correlated to each other, which make sense because people most likely to married a particular age and at the age of 50 most of the peoples are married.
9. Features like: 'Age', 'Hypertension', 'Glucose Level', 'Heart Disease' and 'Ever Married' are the most correlated features with Stroke feature, it make sense because according to our research they all make high impact in stroke cases. (Positive Correlation)
10. Work Type also impact on the stroke cases. (Negative Correlation)

Machine Learning Model:

1. Sampling the dataset to make a equal classes dataset or balance the dataset to make most accurate model.
2. Scaling the training and testing independent data to make all values in range of 0 to 1 using MinMaxScaler() method.
3. Identifying the accuracy and log loss for all machine-learning algorithms. And pick top 3 algorithms.
4. Evaluate top 3 algorithms by identifying classification report and confusion report.
5. So, we got algorithms like: Random Forest, Decision Tree and K Nearest Neighbor according to log loss, accuracy and other evaluation metrices.
6. Random Forest and Decision Tree perform out-standing. So, we will work on it to make our prediction model and check it on new data.
7. 'Age', 'Glucose Level' and 'BMI' has the high feature importance for our models (both Random Forest and Decision Tree).
8. Random Forest is capable to contain whole area in our ROC plot, which means Random Forest is capable to draw accurate prediction on new data.
9. Althrogh, Decision Tree also perform amazing in ROC plot with 0.97 or 97% prediction rate.
10. We will use Decision Tree Model to make predictions.

NOTE: As we know from data analysis part that our dataset have imbalanced classes and later we perform 'Resampling' method to handle it. Although, we successfully done it and our models perform well. But, this also show biases because, in resampling we will take random samples or duplicate sample to balance our dataset, that's why our model is not much perform well for new data or real time data. This is a learning purpose project

which gives you idea how to do analysis, resampling and making machine-learning model. So, don't execute same to same as I do, something try with yourself to handle in a best way.

In []: